

STAT 331 Applied Linear Models

University of Waterloo - Fall 2025

10th November 2025

Instructor: Kun Liang

L^AT_EX: Xing Liu

Contents

1	Simple Linear Regression (SLR)	2
1.1	Introduction	2
1.2	Least Squares Estimation	3
1.3	SLR: Inference	9
1.4	SLR: ANOVA and Cochran's Theorem	12
2	Multiple Linear Regression (MLR)	17
2.1	Linear Algebra Review and MVN Distribution	17
2.2	Introduction to MLR	21
2.3	Least Squares Estimation (LSE)	22
2.4	MLR: Inference	26
2.5	MLR: ANOVA	34
2.6	Geometric Interpretation of LSE	37
3	Model Specification	46
4	Model Checking	55

1 Simple Linear Regression (SLR)

Lecture 1, 2025/09/04

Definition 1.1 (Linear Regression Model).

We say the regression model

$$y = f(x_1, x_2, \dots, x_p) + \epsilon$$

is **linear** if it is linear in the parameters, i.e. $\frac{\partial f}{\partial \beta_i}$ does not depend on β_i for $i = 0, 1, \dots, p$.

Lecture 2, 2025/09/09

1.1 Introduction

Definition 1.2 (SLR Population Model).

Population model.

$$\text{Response: } y = \beta_0 + \beta_1 x + \epsilon,$$

where β_0, β_1 are regression coefficients, x is the predictor, and ϵ is the random error.

Definition 1.3 (SLR Sample Model).

Sample: n pairs of observations $(x_i, y_i), i = 1, 2, \dots, n$.

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i,$$


where β_0 is the intercept and β_1 is the slope.

	Fixed/Random	Known/Unknown
x_i 's	Fixed	Known
β_i 's	Fixed	Unknown
ϵ_i 's	Random	Unknown
y_i 's	Random	Known

SLR Assumptions

- (1) $\mathbb{E}[\epsilon_i] = 0 \implies \mathbb{E}[y_i] = \beta_0 + \beta_1 x_i$.
- (2) $\epsilon_1, \dots, \epsilon_n$ are statistically independent.
- (3) $\text{Var}(\epsilon_i) = \sigma^2 \implies \text{Var}(y_i) = \sigma^2$.
- (4) ϵ_i is normally distributed, i.e. $\epsilon_i \sim N(0, \sigma^2) \implies y_i \sim N(\beta_0 + \beta_1 x_i, \sigma^2)$.

Note. The first three assumptions are called the **Gauss-Markov assumptions**.



The slope β_1 is often of primary interest. We have the following interpretations of β_1 .

- (1) $\beta_1 = \mathbb{E}[y \mid x = a + 1] - \mathbb{E}[y \mid x = a]$.
- (2) If $\beta_1 = 0$, then $\mathbb{E}[y \mid x] = \beta_0$, i.e. y does not depend on x .

1.2 Least Squares Estimation

Task: Given the sample observation (x_i, y_i) , $i = 1, 2, \dots, n$, we want to find (β_0, β_1) such that

$$S(\beta_0, \beta_1) = \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2$$

is minimized.

To minimize the objective function $S(\beta_0, \beta_1)$, we set the partial derivatives to zero:

$$\frac{\partial S(\beta_0, \beta_1)}{\partial \beta_0} = -2 \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i) = 0, \quad (1.1)$$

$$\frac{\partial S(\beta_0, \beta_1)}{\partial \beta_1} = -2 \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i) x_i = 0. \quad (1.2)$$

Then, to solve for β_0, β_1 ,

$$(1.1) \implies n\beta_0 = \sum_{i=1}^n y_i - \beta_1 \sum_{i=1}^n x_i, \text{ so}$$

$$\beta_0 = \bar{y} - \beta_1 \bar{x}, \text{ where } \bar{y} = \frac{1}{n} \sum_{i=1}^n y_i \text{ and } \bar{x} = \frac{1}{n} \sum_{i=1}^n x_i.$$

$$(1.2) \implies \sum_{i=1}^n x_i y_i - \beta_0 \sum_{i=1}^n x_i - \beta_1 \sum_{i=1}^n x_i^2 = 0. \text{ Substitute } \beta_0 \text{ into this, we have}$$

$$\begin{aligned} \sum_{i=1}^n x_i y_i - (\bar{y} - \beta_1 \bar{x}) \sum_{i=1}^n x_i - \beta_1 \sum_{i=1}^n x_i^2 &= 0 \\ \sum_{i=1}^n x_i y_i - \bar{y} \sum_{i=1}^n x_i + \beta_1 \bar{x} \sum_{i=1}^n x_i - \beta_1 \sum_{i=1}^n x_i^2 &= 0 \\ \beta_1 &= \frac{\sum_{i=1}^n x_i y_i - n \bar{x} \bar{y}}{\sum_{i=1}^n x_i^2 - n \bar{x}^2} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} = \frac{S_{xy}}{S_{xx}}, \end{aligned}$$

Lecture 3, 2025/09/11

Recall from STAT 231:

$$\text{Sample variance of } x: S_x^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2,$$

$$\text{Sample covariance: } q_{xy} = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}).$$

Then,

$$\beta_1 = \frac{S_{xy}}{S_{xx}} = \frac{q_{xy}}{S_x^2}.$$

Definition 1.4 (Pearson Correlation Coefficient).

The **Pearson correlation coefficient** between x and y is

$$\rho_{xy} = \frac{\text{Cov}(X, Y)}{\sqrt{\text{Var}(X) \text{Var}(Y)}}.$$

Definition 1.5 (Sample Correlation Coefficient).

The **sample correlation coefficient** between x and y is

$$r_{xy} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2 \sum_{i=1}^n (y_i - \bar{y})^2}}.$$

Then, we have $\beta_1 = r_{xy} \frac{\sqrt{S_{yy}}}{\sqrt{S_{xx}}}$.

Thus, the least squares estimates (LSEs) are:

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}, \quad \hat{\beta}_1 = \frac{S_{xy}}{S_{xx}}.$$

Proposition 1.1 (Properties of $\hat{\beta}_0$ and $\hat{\beta}_1$).

- (1) The least squares estimates (LSEs) are unbiased, i.e. $\mathbb{E}[\hat{\beta}_0] = \beta_0, \mathbb{E}[\hat{\beta}_1] = \beta_1$.
- (2) $\text{Var}(\hat{\beta}_0) = \sigma^2 \left(\frac{1}{n} + \frac{\bar{x}^2}{S_{xx}} \right), \text{Var}(\hat{\beta}_1) = \frac{\sigma^2}{S_{xx}}$.
- (3) $\text{Cov}(\hat{\beta}_0, \hat{\beta}_1) = -\frac{\sigma^2 \bar{x}}{S_{xx}}$.

Proof.

(1) We have

$$\begin{aligned} \mathbb{E}[\hat{\beta}_1] &= \mathbb{E}\left[\frac{S_{xy}}{S_{xx}}\right] = \mathbb{E}\left[\frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{S_{xx}}\right] \\ &= \mathbb{E}\left[\frac{\sum_{i=1}^n (x_i - \bar{x})y_i}{S_{xx}}\right] \end{aligned}$$

Note that $\hat{\beta}_1 = \sum_{i=1}^n c_i y_i$, where $c_i = \frac{x_i - \bar{x}}{S_{xx}}$.

Remark (Properties of c_i 's).

- (i) $\sum_{i=1}^n c_i = 0$.
- (ii) $\sum_{i=1}^n c_i x_i = \frac{1}{S_{xx}} \sum_{i=1}^n (x_i - \bar{x}) x_i = \frac{1}{S_{xx}} (\sum_{i=1}^n x_i^2 - n\bar{x}^2) = 1$.
- (iii) $\sum_{i=1}^n c_i^2 = \frac{1}{S_{xx}^2} \sum_{i=1}^n (x_i - \bar{x})^2 = \frac{1}{S_{xx}}$.

Then,

$$\begin{aligned} \mathbb{E}[\hat{\beta}_1] &= \mathbb{E}\left[\sum_{i=1}^n c_i y_i\right] = \sum_{i=1}^n c_i \mathbb{E}[y_i] \\ &= \sum_{i=1}^n c_i (\beta_0 + \beta_1 x_i) = \beta_0 \sum_{i=1}^n c_i + \beta_1 \sum_{i=1}^n c_i x_i = \beta_1. \end{aligned}$$

Similarly,

$$\begin{aligned} \mathbb{E}[\hat{\beta}_0] &= \mathbb{E}[\bar{y} - \hat{\beta}_1 \bar{x}] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[y_i] - \bar{x} \mathbb{E}[\hat{\beta}_1] \\ &= \frac{1}{n} \sum_{i=1}^n (\beta_0 + \beta_1 x_i) - \bar{x} \beta_1 \\ &= \beta_0 + \beta_1 \bar{x} - \bar{x} \beta_1 = \beta_0. \end{aligned}$$

(2) We have

$$\begin{aligned} \text{Var}(\hat{\beta}_1) &= \text{Var}\left(\sum_{i=1}^n c_i y_i\right) = \sum_{i=1}^n c_i^2 \text{Var}(y_i) \\ &= \sigma^2 \sum_{i=1}^n c_i^2 = \frac{\sigma^2}{S_{xx}}. \end{aligned}$$

Similarly,

$$\begin{aligned}
\text{Var}(\hat{\beta}_0) &= \text{Var}(\bar{y} - \hat{\beta}_1 \bar{x}) = \text{Var}\left(\frac{1}{n} \sum_{i=1}^n y_i - \bar{x} \sum_{i=1}^n c_i y_i\right) \\
&= \text{Var}\left(\sum_{i=1}^n \left(\frac{1}{n} - c_i \bar{x}\right) y_i\right) = \sum_{i=1}^n \left(\frac{1}{n} - c_i \bar{x}\right)^2 \text{Var}(y_i) \\
&= \sigma^2 \sum_{i=1}^n \left(\frac{1}{n} - c_i \bar{x}\right)^2 = \sigma^2 \left(\sum_{i=1}^n \frac{1}{n^2} - 2 \frac{\bar{x}}{n} \sum_{i=1}^n c_i + \bar{x}^2 \sum_{i=1}^n c_i^2\right) \\
&= \sigma^2 \left(\frac{1}{n} + \frac{\bar{x}^2}{S_{xx}}\right).
\end{aligned}$$

(3) We have

$$\begin{aligned}
\text{Cov}(\hat{\beta}_0, \hat{\beta}_1) &= \text{Cov}(\bar{y} - \hat{\beta}_1 \bar{x}, \hat{\beta}_1) \\
&= \text{Cov}\left(\sum_{i=1}^n \left(\frac{1}{n} - c_i \bar{x}\right) y_i, \sum_{i=1}^n c_i y_i\right) \\
&= \sum_{i=1}^n \sum_{j=1}^n \left(\frac{1}{n} - c_i \bar{x}\right) c_j \text{Cov}(y_i, y_j)
\end{aligned}$$

Note that $\text{Cov}(y_i, y_j) = \begin{cases} \sigma^2, & \text{if } i = j, \\ 0, & \text{if } i \neq j. \end{cases}$ Then,

$$\begin{aligned}
\text{Cov}(\hat{\beta}_0, \hat{\beta}_1) &= \sigma^2 \sum_{i=1}^n \left(\frac{1}{n} - c_i \bar{x}\right) c_i \\
&= \sigma^2 \left(\frac{1}{n} \sum_{i=1}^n c_i - \bar{x} \sum_{i=1}^n c_i^2\right) = -\frac{\sigma^2 \bar{x}}{S_{xx}}.
\end{aligned}$$

□

Definition 1.6 (Residual).

The **residual** for the i -th observation is the difference between the observed and the fitted value:

$$r_i = y_i - \hat{y}_i = y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i.$$

Under the least squares fit, we have the following.

Proposition 1.2 (Properties of Residuals).

- (1) $\sum_{i=1}^n r_i = 0.$
- (2) $\sum_{i=1}^n r_i x_i = 0.$
- (3) $\sum_{i=1}^n r_i \hat{y}_i = 0$, where $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i.$

Lecture 4, 2025/09/16

To find the estimator of σ^2 , note that

$$\epsilon_i = y_i - (\beta_0 + \beta_1 x_i)$$

$$r_i = y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i).$$

If ϵ_i 's are known, we can estimate σ^2 by $\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n \epsilon_i^2$ because $\epsilon_i \sim N(0, \sigma^2)$ and

$$\mathbb{E}[\hat{\sigma}^2] = \mathbb{E}[\epsilon_i]^2 + \text{Var}(\epsilon_i) = \sigma^2.$$

However $\mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n r_i^2\right] \neq \sigma^2$. If we define

$$s^2 = \frac{1}{n-2} \sum_{i=1}^n r_i^2,$$

then, $\mathbb{E}[s^2] = \sigma^2$.

Note. The degree of freedom is $n - 2$ because we have estimated two parameters, β_0 and β_1 .

Example (Degree of Freedom Example).

If $y_i \sim N(\mu, \sigma^2)$, then the sample variance estimator is

$$\hat{\sigma}^2 = \frac{1}{n-1} \sum_{i=1}^n (y_i - \hat{\mu})^2,$$

where the degree of freedom is $n - 1$ because we have estimated one parameter. Note that

$$\mathbb{E}[\hat{\sigma}^2] = \sigma^2.$$

1.3 SLR: Inference

Proposition 1.3. Under the assumptions that ϵ_i 's are independent and normally distributed,

$$\hat{\beta}_1 \sim N\left(\beta_1, \frac{\sigma^2}{S_{xx}}\right).$$

Remark. If σ^2 is known, then

$$\frac{\hat{\beta}_1 - \beta_1}{\sqrt{\sigma^2/S_{xx}}} \sim N(0, 1).$$

In most practice, σ^2 is unknown, and we replace σ^2 by s^2 :

$$\frac{\hat{\beta}_1 - \beta_1}{\sqrt{s^2/S_{xx}}} \sim t_{n-2}.$$

Proposition 1.4 (Confidence Interval for β_1).

A $100(1 - \alpha)\%$ confidence interval (CI) for β_1 is

$$[\hat{\beta}_1 - t_{n-2, \alpha/2} \text{SE}(\hat{\beta}_1), \hat{\beta}_1 + t_{n-2, \alpha/2} \text{SE}(\hat{\beta}_1)],$$


or equivalently

$$\hat{\beta}_1 \pm t_{n-2, \alpha/2} \text{SE}(\hat{\beta}_1), \text{ where } \text{SE}(\hat{\beta}_1) = \frac{s}{\sqrt{S_{xx}}}.$$

Proof. Note that

$$\begin{aligned} \mathbb{P}\left(-t_{n-2, \alpha/2} \leq \frac{\hat{\beta}_1 - \beta_1}{\text{SE}(\hat{\beta}_1)} \leq t_{n-2, \alpha/2}\right) &= 1 - \alpha \\ \mathbb{P}(\hat{\beta}_1 - t_{n-2, \alpha/2} \text{SE}(\hat{\beta}_1) \leq \beta_1 \leq \hat{\beta}_1 + t_{n-2, \alpha/2} \text{SE}(\hat{\beta}_1)) &= 1 - \alpha. \end{aligned}$$

□



Hypothesis Testing for β_1

We test

$$H_0 : \beta_1 = \beta_1^* \text{ vs. } H_a : \beta_1 \neq \beta_1^*.$$

Under H_0 , the test statistic is

$$t = \frac{\hat{\beta}_1 - \beta_1^*}{\text{SE}(\hat{\beta}_1)} \sim t_{n-2}.$$

Reject H_0 at significance level α if $|t| > t_{n-2, \alpha/2}$. Alternatively, compute the p -value:

$$p = 2\mathbb{P}(T > |t|) \quad \text{where } T \sim t_{n-2}$$

and reject H_0 if $p < \alpha$. Typically, we are interested in testing $H_0 : \beta_1 = 0$.

Note. $p = \mathbb{P}(|T| \geq |t|) = 2\mathbb{P}(T \geq |t|)$ because t -distribution is symmetric about 0.


Inference for Mean Response

For the predictor value x_0 , the mean response is $\mu_0 = \beta_0 + \beta_1 x_0$. We estimate μ_0 by

$$\hat{\mu}_0 = \hat{\beta}_0 + \hat{\beta}_1 x_0.$$

We can use $\hat{\mu}_0 = \bar{y} - \hat{\beta}_1 \bar{x} + \hat{\beta}_1 x_0 = \bar{y} + \hat{\beta}_1 (x_0 - \bar{x})$ to show that

$$\hat{\mu}_0 = \sum_i d_i y_i, \text{ where } d_i = \frac{1}{n} + \frac{(x_i - \bar{x})(x_0 - \bar{x})}{S_{xx}}.$$



Note. The black line is the true regression line and the blue line is the fitted line based on estimates.

Proposition 1.5 (Expectation & Variance of $\hat{\mu}_0$).

For $\hat{\mu}_0 = \hat{\beta}_0 + \hat{\beta}_1 x_0$,

$$\mathbb{E}[\hat{\mu}_0] = \mu_0, \quad \text{Var}(\hat{\mu}_0) = \left[\frac{1}{n} + \frac{(x_0 - \bar{x})^2}{S_{xx}} \right] \sigma^2.$$

Question: What is your best guess of y given $x = x_p$?

Answer: $\hat{y}_p$, where $\hat{y}_p = \hat{\beta}_0 + \hat{\beta}_1 x_p$.

Note that

$$\mathbb{E}[y_p - \hat{y}_p] = 0, \quad \text{Var}(y_p - \hat{y}_p) = \left[1 + \frac{1}{n} + \frac{(x_p - \bar{x})^2}{S_{xx}} \right] \sigma^2.$$

Since σ^2 is unknown, we use s^2 and get

$$\frac{y_p - \hat{y}_p}{\text{SE}(y_p - \hat{y}_p)} \sim t_{n-2}, \quad \text{where } \text{SE}(y_p - \hat{y}_p) = s \sqrt{1 + \frac{1}{n} + \frac{(x_p - \bar{x})^2}{S_{xx}}}.$$

Note. $y_p = \beta_0 + \beta_1 x_p + \epsilon_p$ is a new observation at $x = x_p$.

Proposition 1.6 (Prediction Interval for y_p).

A $100(1 - \alpha)\%$ prediction interval (PI) for y_p is


$$\hat{y}_p \pm t_{n-2, \alpha/2} \text{SE}(y_p - \hat{y}_p), \quad \text{where } \text{SE}(y_p - \hat{y}_p) = s \sqrt{1 + \frac{1}{n} + \frac{(x_p - \bar{x})^2}{S_{xx}}}.$$

1.4 SLR: ANOVA and Cochran's Theorem

Analysis of Variance (ANOVA): How well does our regression model fit the response variable?

Total variation in y_i 's is measured by the total sum of squares:

$$SST = \sum_{i=1}^n (y_i - \bar{y})^2.$$



Proposition 1.7 (ANOVA Decomposition).

We have the following decomposition.

$$SST = SSR + SSE,$$

where

$$SSR = \sum_{i=1}^n (\hat{y}_i - \bar{y})^2 \text{ is the regression sum of squares,}$$

$$SSE = \sum_{i=1}^n r_i^2 = \sum_{i=1}^n (y_i - \hat{y}_i)^2 \text{ is the residual sum of squares.}$$

Proof. Notice that $y_i - \bar{y} = y_i - \hat{y}_i + \hat{y}_i - \bar{y}$. Now, we have

$$\begin{aligned}
\text{SST} &= \sum_{i=1}^n (y_i - \hat{y}_i + \hat{y}_i - \bar{y})^2 \\
&= \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \sum_{i=1}^n (\hat{y}_i - \bar{y})^2 + 2 \sum_{i=1}^n (y_i - \hat{y}_i)(\hat{y}_i - \bar{y}) \\
&= \text{SSE} + \text{SSR} + 2 \sum_{i=1}^n r_i (\hat{y}_i - \bar{y}) \\
&= \text{SSE} + \text{SSR}.
\end{aligned}$$

□

Note. SSR is the variability explained by the model; SSE is the variability not explained by the model.

ANOVA for Testing $H_0 : \beta_1 = 0$

If H_0 is true, $y_i = \beta_0 + \epsilon_i$. SSR should be “small” relatively to SSE.

Proposition 1.8. Under $H_0 : \beta_1 = 0$, we have the following.

- (1) $\frac{\text{SSR}}{\sigma^2} \sim \chi_1^2$.
- (2) $\frac{\text{SST}}{\sigma^2} \sim \chi_{n-1}^2$.
- (3) $\frac{\text{SSE}}{\sigma^2} \sim \chi_{n-2}^2$, and SSE is independent of SSR.

Proof.

(1) We have

$$\begin{aligned}
\text{SSR} &= \sum_{i=1}^n (\hat{y}_i - \bar{y})^2 = \sum_{i=1}^n (\hat{\beta}_0 + \hat{\beta}_1 x_i - \bar{y})^2 \\
&= \sum_{i=1}^n (\bar{y} - \hat{\beta}_1 \bar{x} + \hat{\beta}_1 x_i - \bar{y})^2 \\
&= \hat{\beta}_1^2 \sum_{i=1}^n (x_i - \bar{x})^2 = \hat{\beta}_1^2 S_{xx}.
\end{aligned}$$

Recall that $\hat{\beta}_1 \sim N\left(\beta_1, \frac{\sigma^2}{S_{xx}}\right)$. Then, we have $\left(\frac{\hat{\beta}_1 - \beta_1}{\sigma/\sqrt{S_{xx}}}\right)^2 \sim \chi_1^2$. Under $H_0 : \beta_1 = 0$, we have

$$\frac{\hat{\beta}_1^2 S_{xx}}{\sigma^2} \sim \chi_1^2 \implies \frac{\text{SSR}}{\sigma^2} \sim \chi_1^2.$$

(2) Under $H_0 : \beta_1 = 0$, we have $y_1, \dots, y_n \stackrel{\text{iid}}{\sim} N(\beta_0, \sigma^2)$. That is, $\frac{y_i - \beta_0}{\sigma} \sim N(0, 1)$. Thus,

$$\sum_{i=1}^n \left(\frac{y_i - \beta_0}{\sigma} \right)^2 \sim \chi_n^2.$$

Furthermore, $\left(\frac{\bar{y} - \beta_0}{\sigma/\sqrt{n}} \right)^2 \sim \chi_1^2$ because $\bar{y} \sim N\left(\beta_0, \frac{\sigma^2}{n}\right)$. Now,

$$\begin{aligned} \text{SST} &= \sum_{i=1}^n (y_i - \bar{y})^2 \\ &= \sum_{i=1}^n [(y_i - \beta_0) - (\bar{y} - \beta_0)]^2 \\ &= \sum_{i=1}^n (y_i - \beta_0)^2 - n(\bar{y} - \beta_0)^2. \end{aligned}$$

Thus, we have

$$\frac{\sum_{i=1}^n (y_i - \beta_0)^2}{\sigma^2} = \frac{\text{SST}}{\sigma^2} + \left(\frac{\bar{y} - \beta_0}{\sigma/\sqrt{n}} \right)^2 \iff S = Q_1 + Q_2,$$

where $S \sim \chi_n^2$. For $Q_1 = \frac{\text{SST}}{\sigma^2}$, there are only $n - 1$ degrees of freedom because $\sum_{i=1}^n (y_i - \bar{y}) = 0$.

For $Q_2 = \left(\frac{\bar{y} - \beta_0}{\sigma/\sqrt{n}} \right)^2$, we only have $\bar{y}$ and so $d_2 = 1$. By Cochran's theorem, Q_1 and Q_2 are independent, and $\frac{\text{SST}}{\sigma^2} \sim \chi_{n-1}^2$.

(3) We have

$$\begin{aligned} \frac{\text{SST}}{\sigma^2} &= \frac{\text{SSE}}{\sigma^2} + \frac{\text{SSR}}{\sigma^2} \\ \underbrace{\frac{\text{SST}}{\sigma^2}}_{\sim \chi_{n-1}^2} &= Q_1 + \underbrace{\frac{\text{SSR}}{\sigma^2}}_{\sim \chi_1^2}. \end{aligned}$$

The degrees of freedom of $\text{SSE} = \sum_{i=1}^n r_i^2$ is $n - 2$ because $\sum_{i=1}^n r_i = 0$ and $\sum_{i=1}^n r_i x_i = 0$. By Cochran's theorem, $\frac{\text{SSE}}{\sigma^2} \sim \chi_{n-2}^2$ and SSE is independent of SSR.

□

Theorem 1.9 (Cochran's Theorem).

Suppose that $U_i \stackrel{\text{iid}}{\sim} N(0, 1)$ for $i = 1, 2, \dots, n$, and

$$\sum_{i=1}^n U_i^2 = Q_1 + Q_2.$$

Let d_1 and d_2 be the degrees of freedom of Q_1 and Q_2 , i.e. the numbers of linearly independent linear combinations of the U_i 's in Q_1 and Q_2 . If $d_1 + d_2 = n$, then Q_1 and Q_2 are independent, and furthermore,

$$Q_1 \sim \chi_{d_1}^2, \quad Q_2 \sim \chi_{d_2}^2.$$


Proposition 1.8 above allows us to define the following F -statistic (with $\beta_1 = 0$ under H_0):

$$F = \frac{\frac{\text{SSR}}{\sigma^2}/1}{\frac{\text{SSE}}{\sigma^2}/(n-2)} = \frac{\text{SSR}}{\text{SSE}/(n-2)} = \frac{\text{MSR}}{\text{MSE}} \sim F(1, n-2).$$

Note that MSR is the mean squares of regression and MSE is the mean squares of error. To test $H_0 : \beta_1 = 0$ vs. $H_a : \beta_1 \neq 0$, we reject H_0 at the level α if

$$F > F_\alpha(1, n-2), \text{ where } F_\alpha(1, n-2) \text{ is the } 100(1-\alpha)\% \text{-tile of } F(1, n-2).$$

The F -distribution looks like this:



Lecture 6, 2025/09/23

To test $H_0 : \beta_1 = 0$, we use the F -statistic:

$$F = \frac{SSR}{SSE/(n-2)} = \frac{MSR}{MSE} \sim F(1, n-2).$$

Recall that

$$\frac{\hat{\beta}_1 - \beta_1^*}{SE(\hat{\beta}_1)} \sim t_{n-2}$$

and we can use this to test $H_0 : \beta_1 = \beta_1^*$. In SLR, testing $H_0 : \beta_1 = 0$ using t -test and F -test are equivalent. If $X \sim t_{n-2}$, then $X^2 \sim F(1, n-2)$.

Source of Variation	Sum of Squares	Degrees of Freedom	Mean Squares	F
Regression	SSR	1	$MSR = \frac{SSR}{1}$	$\frac{MSR}{MSE}$
Residual	SSE	$n-2$	$MSE = \frac{SSE}{n-2}$	
Total	SST	$n-1$		

Definition 1.7 (Coefficient of Determination).

The **coefficient of determination** is defined as

$$R^2 = \frac{SSR}{SST}.$$

It measures the proportion of total variation in the response variable y around the mean that is explained by the regression model.

$0 \leq R^2 \leq 1$, $R^2 = 0$ means none of the variability in y is explained by x . In SLR, we have

$$R^2 = \frac{SSR}{SST} = \frac{\hat{\beta}_1^2 S_{xx}}{S_{yy}} = \frac{\left(\frac{S_{xy}}{S_{xx}}\right)^2 S_{xx}}{S_{yy}} = \frac{S_{xy}^2}{S_{yy} S_{xx}}.$$

Recall that the sample correlation coefficient is

$$r = \frac{S_{xy}}{\sqrt{S_{xx} S_{yy}}}$$

and thus $R^2 = r^2$.

2 Multiple Linear Regression (MLR)

2.1 Linear Algebra Review and MVN Distribution

Definition 2.1 (Random Vectors).

Let $\mathbf{Y} = (y_1, \dots, y_n)^\top$ be a vector of random variables. Then, $\mathbf{Y}$ is called a **random vector**.

Note that

$$\mathbb{E}[y_i] = \mu_i, \quad \text{Var}(y_i) = \sigma_i^2, \quad \text{Cov}(y_i, y_j) = \sigma_{ij}.$$

Also,

$$\mathbb{E}[\mathbf{Y}] = \begin{bmatrix} \mu_1 \\ \mu_2 \\ \vdots \\ \mu_n \end{bmatrix}, \quad \text{Var}(\mathbf{Y}) = \text{Cov}(\mathbf{Y}, \mathbf{Y}) = \begin{bmatrix} \sigma_1^2 & \sigma_{12} & \cdots & \sigma_{1n} \\ \sigma_{21} & \sigma_2^2 & \cdots & \sigma_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ \sigma_{n1} & \sigma_{n2} & \cdots & \sigma_n^2 \end{bmatrix}.$$

Remark. Note that

$$\begin{aligned} \text{Var}(\mathbf{Y}) &= \mathbb{E}[(\mathbf{Y} - \mathbb{E}[\mathbf{Y}])(\mathbf{Y} - \mathbb{E}[\mathbf{Y}])^\top] \\ &= \begin{bmatrix} \text{Var}(y_1) & \text{Cov}(y_1, y_2) & \cdots & \text{Cov}(y_1, y_n) \\ \text{Cov}(y_2, y_1) & \text{Var}(y_2) & \cdots & \text{Cov}(y_2, y_n) \\ \vdots & \vdots & \ddots & \vdots \\ \text{Cov}(y_n, y_1) & \text{Cov}(y_n, y_2) & \cdots & \text{Var}(y_n) \end{bmatrix} \end{aligned}$$

Example. If $y_1, \dots, y_n$ are independent, $\text{Cov}(y_i, y_j) = 0 \forall i \neq j$. Furthermore, $\text{Var}(y_i) = \sigma^2 \forall i$. This implies that

$$\text{Var}(\mathbf{Y}) = \sigma^2 I, \quad \text{where } I = \begin{bmatrix} 1 & 0 & \cdots & 0 \\ 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 1 \end{bmatrix} \text{ is the identity matrix.}$$

Proposition 2.1 (Basic Properties).

Suppose $A = (a_{ij})_{m \times n}$, $\mathbf{b} = (b_1, \dots, b_m)^\top$, $\mathbf{c} = (c_1, \dots, c_n)^\top$ are matrix and vectors of constants. Then,

- (1) $\mathbb{E}[A\mathbf{Y} + \mathbf{b}] = A\mathbb{E}[\mathbf{Y}] + \mathbf{b}$.
- (2) $\text{Var}(\mathbf{Y} + \mathbf{c}) = \text{Var}(\mathbf{Y})$.
- (3) $\text{Var}(A\mathbf{Y}) = A \text{Var}(\mathbf{Y}) A^\top$.
- (4) $\text{Var}(A\mathbf{Y} + \mathbf{b}) = A \text{Var}(\mathbf{Y}) A^\top$.

Definition 2.2 (Differentiation: Linear & Quadratic Forms).

Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a function and $\mathbf{Y}$ be a vector. Then,

- (1) If $f = f(\mathbf{Y})$, then

$$\frac{df}{d\mathbf{Y}} = \begin{bmatrix} \frac{\partial f}{\partial y_1} \\ \vdots \\ \frac{\partial f}{\partial y_n} \end{bmatrix}.$$

- (2) If $f = \mathbf{c}^\top \mathbf{Y} = \sum c_i y_i$, then

$$\frac{df}{d\mathbf{Y}} = \mathbf{c}.$$

- (3) If $A = A^\top$ and $f = \mathbf{Y}^\top A \mathbf{Y} = \sum_{i=1}^n \sum_{j=1}^n a_{ij} y_i y_j$, then

$$\frac{df}{d\mathbf{Y}} = 2A\mathbf{Y}.$$

In this case, f is called a **quadratic form**.

Definition 2.3 (Trace).

The **trace** of a square matrix $A = (a_{ij})_{n \times n}$ is defined as

$$\text{tr}(A) = \sum_{i=1}^n a_{ii}.$$

Remark. We have $\text{tr}(BC) = \text{tr}(CB)$.

Definition 2.4 (Rank).

The **rank** of a matrix A is defined as

$\text{rank}(A)$ = the maximum number of linearly independent columns (or rows) of A .

Definition 2.5 (Linearly Independence).

Random vectors $\mathbf{Y}_1, \dots, \mathbf{Y}_n$ are linearly independent $\iff$

$$c_1 \mathbf{Y}_1 + \dots + c_n \mathbf{Y}_n = \mathbf{0} \implies c_1 = c_2 = \dots = c_n = 0.$$

Definition 2.6 (Orthogonality, Orthogonal Matrix).

Random vectors $\mathbf{X}, \mathbf{Y}$ are **orthogonal** if $\mathbf{X}^\top \mathbf{Y} = 0$.

A square matrix A is an **orthogonal matrix** if $A^\top A = AA^\top = I$.

Definition 2.7 (Eigenvalues, Eigenvectors).

Let A be a square matrix. A scalar λ is an **eigenvalue** of A if $\exists$ a non-zero vector $\mathbf{v}$ such that

$$A\mathbf{v} = \lambda\mathbf{v}.$$

The vector $\mathbf{v}$ is called an **eigenvector** corresponding to the eigenvalue λ .

Definition 2.8 (Idempotent Matrix).

A square matrix A is **idempotent** if $A^2 = A$.

Proposition 2.2. Let A be an idempotent matrix.

- (1) All eigenvalues of A are either 0 or 1.
- (2) If A is symmetric (i.e. $A = A^\top$), then $\exists$ orthogonal P s.t.

$$A = P\Lambda P^\top, \quad \text{where } \Lambda = \text{diag}(1, \dots, 1, 0, \dots, 0).$$

Then, $\text{tr}(A) = \text{rank}(A) = \#$ of eigenvalues equal to 1.

Proof.

- (1) Let λ be an eigenvalue of A with corresponding eigenvector $\mathbf{v}$. Then,

$$A^2\mathbf{v} = A(A\mathbf{v}) = A(\lambda\mathbf{v}) = \lambda A\mathbf{v} = \lambda^2\mathbf{v}.$$

But since A is idempotent, $A^2\mathbf{v} = A\mathbf{v} = \lambda\mathbf{v}$. Thus,

$$\lambda^2\mathbf{v} = \lambda\mathbf{v} \implies \lambda(\lambda - 1)\mathbf{v} = 0.$$

- (2) For the $\text{tr}(A) = \text{rank}(A)$ part, note that

$$\text{tr}(A) = \text{tr}(P\Lambda P^\top) = \text{tr}(\Lambda P^\top P) = \text{tr}(\Lambda) = \text{rank}(A).$$

□

Definition 2.9 (Multivariate Normal Distribution (MVN)).

$\mathbf{Y} \sim \text{MVN}(\boldsymbol{\mu}, \Sigma)$ if $\mathbf{Y}$ has the following probability density function (PDF):

$$f(\mathbf{y}) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{y} - \boldsymbol{\mu})^\top \Sigma^{-1}(\mathbf{y} - \boldsymbol{\mu})\right).$$

Also, $\mathbb{E}[\mathbf{Y}] = \boldsymbol{\mu}$ and $\text{Var}(\mathbf{Y}) = \Sigma$.

Proposition 2.3 (Properties of MVN).

Suppose $\mathbf{Y}$ is a random vector and A, B are matrices. Then,

- (1) $y_1, \dots, y_n$ are independent $\iff \Sigma$ is diagonal.
- (2) Marginal normality: $y_i \sim N(\mu_i, \sigma_{ii})$.
- (3) If $\mathbf{Y} \sim \text{MVN}(\boldsymbol{\mu}, \Sigma)$ and $Z = A\mathbf{Y}$, then

$$Z \sim \text{MVN}(A\boldsymbol{\mu}, A\Sigma A^\top).$$

- (4) If $\mathbf{Y} \sim \text{MVN}(\vec{0}, \sigma^2 I)$, then

$$\frac{1}{\sigma^2} \mathbf{Y}^\top \mathbf{Y} \sim \chi_n^2.$$

(5) If $U = A\mathbf{Y}$ and $W = B\mathbf{Y}$, then

$$\text{Cov}(U, W) = A\Sigma B^T \quad \text{i.e.} \quad \text{Cov}(A\mathbf{Y}, B\mathbf{Y}) = A \text{Var}(\mathbf{Y}) B^T.$$

Then, $U \perp W \iff A\Sigma B^T = 0$.

Remark. For (5), note that

$$\text{Cov}(U, W) = \mathbb{E}[(U - \mathbb{E}[U])(W - \mathbb{E}[W])^T] = A\mathbb{E}[(\mathbf{Y} - \mathbb{E}[\mathbf{Y}](\mathbf{Y} - \mathbb{E}[\mathbf{Y}])^T]B^T = A\Sigma B^T.$$

Lecture 7, 2025/09/25

2.2 Introduction to MLR

Definition 2.10 (Multiple Linear Regression (MLR)).

The **MLR** model is

$$y_i = \beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip} + \epsilon_i, \quad i = 1, \dots, n,$$

Explanation of the regression coefficients $\beta_1, \dots, \beta_p$:

- β_j : the average amount of increase (decrease) in response when the j th predictor x_j increases (decreases) by 1 unit while holding all other predictors fixed/constant.

The null hypothesis $H_0 : \beta_j = 0$ means that x_j is not (linearly) related to y given all other predictors in the model.

Matrix Form Representation

We can write the n simultaneous linear models as

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}_{n \times 1} = \begin{bmatrix} 1 & x_{11} & x_{12} & \cdots & x_{1p} \\ 1 & x_{21} & x_{22} & \cdots & x_{2p} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n1} & x_{n2} & \cdots & x_{np} \end{bmatrix}_{n \times (p+1)} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_p \end{bmatrix}_{(p+1) \times 1} + \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_n \end{bmatrix}_{n \times 1}$$

Equivalently, we can write

$$\mathbf{Y} = X\boldsymbol{\beta} + \boldsymbol{\epsilon}.$$

And

$$\boldsymbol{\epsilon} \sim \text{MVN}(\vec{0}, \sigma^2 I), \quad \mathbf{Y} \sim \text{MVN}(X\boldsymbol{\beta}, \sigma^2 I).$$

2.3 Least Squares Estimation (LSE)

Least Squares Estimation (LSE) of $\boldsymbol{\beta} = (\beta_0, \beta_1, \dots, \beta_p)^\top$

Once we collect data and observe y_i 's, we minimize

$$\begin{aligned} S(\boldsymbol{\beta}) &= \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_{i1} - \dots - \beta_p x_{ip})^2 \\ &= (\mathbf{Y} - X\boldsymbol{\beta})^\top (\mathbf{Y} - X\boldsymbol{\beta}) \\ &= \mathbf{Y}^\top \mathbf{Y} - \underbrace{\mathbf{Y}^\top X}_{1 \times n} \boldsymbol{\beta} - \boldsymbol{\beta}^\top X^\top \mathbf{Y} + \boldsymbol{\beta}^\top X^\top X \boldsymbol{\beta}. \end{aligned}$$

Note that

$$\mathbf{Y}^\top X \boldsymbol{\beta} = (\mathbf{Y}^\top X \boldsymbol{\beta})^\top = \boldsymbol{\beta}^\top X^\top \mathbf{Y} \quad \text{is a scalar.}$$

Thus,

$$S(\boldsymbol{\beta}) = \mathbf{Y}^\top \mathbf{Y} - 2 \underbrace{\boldsymbol{\beta}^\top X^\top \mathbf{Y}}_{\text{linear}} + \underbrace{\boldsymbol{\beta}^\top X^\top X \boldsymbol{\beta}}_{\text{quadratic form}}.$$

Recall that

$$\frac{d\boldsymbol{\beta}^\top \mathbf{Y}}{d\mathbf{Y}} = \boldsymbol{\beta}, \quad \frac{d\mathbf{Y}^\top A \mathbf{Y}}{d\mathbf{Y}} = 2A\mathbf{Y}.$$

Taking derivative w.r.t. $\boldsymbol{\beta}$, we have

$$\frac{dS(\boldsymbol{\beta})}{d\boldsymbol{\beta}} = -2X^\top \mathbf{Y} + 2X^\top X \boldsymbol{\beta} \stackrel{\text{set}}{=} \vec{0} \implies \hat{\boldsymbol{\beta}} = (X^\top X)^{-1} X^\top \mathbf{Y}.$$

Note. We assume that $(X^\top X)$ is invertible, i.e. $\text{rank}(X) = p + 1$.

Proposition 2.4 (Properties of LSE $\hat{\boldsymbol{\beta}}$).

- (1) $\hat{\boldsymbol{\beta}}$ is an unbiased estimator of $\boldsymbol{\beta}$, i.e. $\mathbb{E}[\hat{\boldsymbol{\beta}}] = \boldsymbol{\beta}$.
- (2) $\text{Var}(\hat{\boldsymbol{\beta}}) = \sigma^2 (X^\top X)^{-1}$.

Proof.

(1) We have

$$\mathbb{E}[\hat{\beta}] = \mathbb{E}[(X^T X)^{-1} X^T \mathbf{Y}] = (X^T X)^{-1} X^T \mathbb{E}[\mathbf{Y}] = (X^T X)^{-1} X^T (X\beta) = \beta.$$

(2) We have

$$\begin{aligned} \text{Var}(\hat{\beta}) &= \text{Var}((X^T X)^{-1} X^T \mathbf{Y}) \\ &= (X^T X)^{-1} X^T \text{Var}(\mathbf{Y}) X [(X^T X)^{-1}]^T \quad \text{since } \text{Var}(A\mathbf{Y}) = A \text{Var}(\mathbf{Y}) A^T \end{aligned}$$

Note that $[(X^T X)^{-1}]^T = [(X^T X)^T]^{-1} = (X^T X)^{-1}$, see remark below.

$$\begin{aligned} &= (X^T X)^{-1} X^T (\sigma^2 I) X (X^T X)^{-1} \\ &= \sigma^2 (X^T X)^{-1} X^T X (X^T X)^{-1} \\ &= \sigma^2 (X^T X)^{-1}. \end{aligned}$$

Remark. In general, $A^{-1}A = I \implies (A^{-1}A)^T = A^T(A^{-1})^T = I^T = I \implies (A^{-1})^T = (A^T)^{-1}$.

□

Proposition 2.5.

(1) Fitted values:

$$\hat{\mathbf{Y}} = X\hat{\beta} = \underbrace{X(X^T X)^{-1} X^T}_{\hat{H}} \mathbf{Y} = H\mathbf{Y}, \quad \text{where } H \text{ is the hat/projection matrix.}$$

(2) H is idempotent and symmetric.

Proof.

(1) Simply substitute the LSE $\hat{\beta} = (X^T X)^{-1} X^T \mathbf{Y}$.

(2) Note that

$$HH = X(X^T X)^{-1} \underbrace{X^T X(X^T X)^{-1}}_I X^T = X(X^T X)^{-1} X^T = H.$$

□

Definition 2.11 (Residuals).

The **residuals** in MLR are defined as

$$\mathbf{r} = \mathbf{Y} - \hat{\mathbf{Y}} = \begin{bmatrix} y_1 - \hat{y}_1 \\ y_2 - \hat{y}_2 \\ \vdots \\ y_n - \hat{y}_n \end{bmatrix} = \begin{bmatrix} r_1 \\ r_2 \\ \vdots \\ r_n \end{bmatrix}.$$

Note that $\mathbf{r} = \mathbf{Y} - H\mathbf{Y} = (I - H)\mathbf{Y}$ where $I - H$ is also idempotent and symmetric.

Proposition 2.6 (Properties of Residuals).

- (1) $\sum_{i=1}^n r_i = 0$.
- (2) $X^T \mathbf{r} = \vec{0}$.
- (3) $\hat{\mathbf{Y}}^T \mathbf{r} = 0$.
- (4) $\mathbb{E}[\mathbf{r}] = \vec{0}$.
- (5) $\text{Var}(\mathbf{r}) = \sigma^2(I - H)$.

Proof.

(1) Using (2), since the first row of X^T consists of all 1's, so $X^T \mathbf{r} = \vec{0}$ implies $\sum r_i = 0$.

(2) Note that

$$\begin{aligned} X^T \mathbf{r} &= X^T (\mathbf{Y} - \hat{\mathbf{Y}}) \\ &= X^T \mathbf{Y} - X^T H \mathbf{Y} \\ &= X^T \mathbf{Y} - \underbrace{X^T X (X^T X)^{-1} X^T}_{I} \mathbf{Y} \\ &= X^T \mathbf{Y} - X^T \mathbf{Y} = \vec{0}_{(p+1) \times 1}. \end{aligned}$$

(3) Note that

$$\hat{\mathbf{Y}}^T \mathbf{r} = (X\hat{\boldsymbol{\beta}})^T \mathbf{r} = \hat{\boldsymbol{\beta}}^T X^T \mathbf{r} = \hat{\boldsymbol{\beta}}^T \vec{0} = 0.$$

(4) Note that

$$\begin{aligned}
\mathbb{E}[\mathbf{r}] &= \mathbb{E}[(I - H)\mathbf{Y}] \\
&= (I - H)\mathbb{E}[\mathbf{Y}] \\
&= (I - X(X^\top X)^{-1}X^\top)X\boldsymbol{\beta} \\
&= X\boldsymbol{\beta} - X \underbrace{(X^\top X)^{-1}X^\top X}_I \boldsymbol{\beta} = \vec{0}_{n \times 1}.
\end{aligned}$$

(5) Note that

$$\begin{aligned}
\text{Var}(\mathbf{r}) &= \text{Var}((I - H)\mathbf{Y}) \\
&= (I - H)\text{Var}(\mathbf{Y})(I - H)^\top \\
&= (I - H)(\sigma^2 I)(I - H) \quad \text{since } I - H \text{ is symmetric} \\
&= \sigma^2(I - H)(I - H) \quad \text{since } I - H \text{ is idempotent} \\
&= \sigma^2(I - H).
\end{aligned}$$

□

An Estimator of σ^2

Proposition 2.7. We have that

$$\hat{\sigma}^2 = \frac{\sum_{i=1}^n r_i^2}{n - (p + 1)}$$

is an unbiased estimator of σ^2 , i.e. $\mathbb{E}[\hat{\sigma}^2] = \sigma^2$.

Note.

(1) In SLR, $p = 1$, we had $\hat{\sigma}^2 = \frac{\sum r_i^2}{n-2}$ where $n - 2$ is the degrees of freedom.

(2) Why $n - (p + 1)$?

$p + 1$ is the # of parameters need to be estimated.

We lose $p + 1$ degrees of freedom because we estimate $p + 1$ parameters $\beta_0, \dots, \beta_p$. The constraints are $X^\top \mathbf{r} = \vec{0}$ which gives $\sum r_i = 0$ and $\sum_{i=1}^n x_{ij}r_i = 0$ for $j = 1, \dots, p$.

Proof. We have

$$\begin{aligned}
\mathbb{E} \left[\sum_{i=1}^n r_i^2 \right] &= \mathbb{E} [\mathbf{r}^\top \mathbf{r}] = \mathbb{E} [\text{tr}(\mathbf{r}^\top \mathbf{r})] \\
&= \mathbb{E} [\text{tr}(\mathbf{r} \mathbf{r}^\top)] \\
&= \text{tr}(\text{Var}(\mathbf{r})) && \text{since } \mathbb{E}[\mathbf{r}] = \vec{0} \\
&= \text{tr}(\sigma^2(I - H)) \\
&= \sigma^2 [\text{tr}(I_{n \times n}) - \text{tr}(H)] \\
&= \sigma^2 [n - \text{tr}(X(X^\top X)^{-1} X^\top)] \\
&= \sigma^2 [n - \text{tr}((X^\top X)^{-1} X^\top X)] && \text{since } \text{tr}(BC) = \text{tr}(CB) \\
&= \sigma^2 [n - \text{tr}(I_{(p+1) \times (p+1)})] \\
&= \sigma^2 [n - (p + 1)].
\end{aligned}$$

□

2.4 MLR: Inference

Sampling distribution of $\hat{\beta}$ and $\hat{\sigma}^2$ under normality assumption $\mathbf{Y} \sim \text{MVN}(X\beta, \sigma^2 I)$.

Proposition 2.8.

- (1) $\hat{\beta} \sim \text{MVN}(\beta, \sigma^2(X^\top X)^{-1})$.
- (2) $\hat{\beta}$ and $\hat{\sigma}^2$ are independent.
- (3) $\frac{(n-(p+1))\hat{\sigma}^2}{\sigma^2} \sim \chi_{n-p-1}^2$.

Proof.

- (1) Recall that $\hat{\beta} = (X^\top X)^{-1} X^\top \mathbf{Y}$. We found the expectation and variance. It follows MVN because $\mathbf{Y} \sim \text{MVN}(X\beta, \sigma^2 I)$ and $\hat{\beta}$ is a linear transformation of $\mathbf{Y}$.

(2) Note that $\hat{\beta} \sim \text{MVN}(\beta, \sigma^2(X^\top X)^{-1})$ from (1) and $\mathbf{r} = (I - H)\mathbf{Y} \sim \text{MVN}(\vec{0}, \sigma^2(I - H))$. Then,

$$\begin{aligned}
\text{Cov}(\hat{\beta}, \mathbf{r}) &= \text{Cov}((X^\top X)^{-1}X^\top \mathbf{Y}, (I - H)\mathbf{Y}) \\
&= (X^\top X)^{-1}X^\top \text{Var}(\mathbf{Y})(I - H)^\top \\
&= (X^\top X)^{-1}X^\top (\sigma^2 I)(I - X(X^\top X)^{-1}X^\top) \\
&= \sigma^2 \left[(X^\top X)^{-1}X^\top - (X^\top X)^{-1} \underbrace{X^\top X (X^\top X)^{-1}X^\top}_I \right] \\
&= \sigma^2 [(X^\top X)^{-1}X^\top - (X^\top X)^{-1}X^\top] \\
&= \mathcal{O}_{(p+1) \times n}.
\end{aligned}$$

Since $\hat{\beta}$ and $\mathbf{r}$ are jointly normal, $\text{Cov}(\hat{\beta}, \mathbf{r}) = \mathcal{O}$ implies that $\hat{\beta} \perp \mathbf{r}$. Since $\hat{\sigma}^2$ is a function of $\mathbf{r}$, we have $\hat{\beta} \perp \hat{\sigma}^2$.

Lecture 9, 2025/10/02

(3) Note that

$$\begin{aligned}
\frac{(n - (p + 1))\hat{\sigma}^2}{\sigma^2} &= \frac{\mathbf{r}^\top \mathbf{r}}{\sigma^2} \\
&= \left(\frac{\mathbf{r}}{\sigma}\right)^\top \left(\frac{\mathbf{r}}{\sigma}\right) \\
&= \mathbf{r}^{*\top} \mathbf{r}^* \quad \text{where } \mathbf{r}^* = \frac{\mathbf{r}}{\sigma} = \frac{(I - H)\mathbf{Y}}{\sigma} \sim \text{MVN}(\vec{0}, I - H)
\end{aligned}$$

By the Eigen Decomposition of a symmetric and idempotent matrix, $\exists$ orthogonal P such that $I - H = P\Lambda P^\top$ where $\Lambda = \text{diag}(1, \dots, 1, 0, \dots, 0)$ with # of 1's = $\text{tr}(I - H) = n - (p + 1)$. Now define a new random vector $\mathbf{z} = P^\top \mathbf{r}^*$. Then,

$$\mathbf{z} \sim \text{MVN}(\vec{0}, P^\top P \Lambda P^\top P) = \text{MVN}(\vec{0}, \Lambda) \implies z_1, \dots, z_{n-p-1} \stackrel{iid}{\sim} N(0, 1) \text{ and } z_{n-p}, \dots, z_n = 0.$$

Thus,


$$\frac{(n - (p + 1))\hat{\sigma}^2}{\sigma^2} = (P\mathbf{z})^\top (P\mathbf{z}) = \mathbf{z}^\top P^\top P \mathbf{z} = \mathbf{z}^\top \mathbf{z} = \sum_{i=1}^{n-p-1} z_i^2 \sim \chi_{n-p-1}^2.$$

Note. Since P is orthogonal, $P^\top P = I$.

□

Example. Let $\mathbf{r} = \begin{bmatrix} r_1 \\ r_2 \end{bmatrix}$ with $r_1 + r_2 = 0$. Let

$$\mathbf{r}^* = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix} \mathbf{r} = \frac{1}{\sqrt{2}} \begin{bmatrix} r_1 + r_2 \\ r_1 - r_2 \end{bmatrix} = \begin{bmatrix} 0 \\ \sqrt{2}r_1 \end{bmatrix}.$$



This example shows some intuition for the proof of Proposition 2.8 (3).

Inference of a single β_i

From Proposition 2.8 (1), we have

$$\hat{\beta}_i \sim N(\beta_i, \sigma^2 \nu_{ii})$$

where ν_{ii} is the i th diagonal element of $(X^\top X)^{-1}$.

From Proposition 2.8 (2) and (3), we have

$$T = \frac{\frac{\hat{\beta}_i - \beta_i}{\sqrt{\sigma^2 \nu_{ii}}}}{\sqrt{\frac{(n - (p + 1))\hat{\sigma}^2 / \sigma^2}{n - (p + 1)}}} \sim t_{n-p-1}$$

since $\frac{\hat{\beta}_i - \beta_i}{\sqrt{\hat{\sigma}^2 \nu_{ii}}} \sim N(0, 1)$ and $(n - (p + 1))\hat{\sigma}^2 / \sigma^2 \sim \chi_{n-p-1}^2$.

We can further simplify T as

$$T = \frac{\hat{\beta}_i - \beta_i}{\sqrt{\hat{\sigma}^2 \nu_{ii}}} \sim t_{n-p-1}$$

and $\text{SE}(\hat{\beta}_i) = \sqrt{\hat{\sigma}^2 \nu_{ii}}$.

Inference of a linear combination of β_i 's

Let $\mathbf{a} = (a_0, a_1, \dots, a_p)^\top$ be a $(p+1)$ -dimensional vector of constants. We are interested in the inference of

$$\theta = \mathbf{a}^\top \boldsymbol{\beta} = \sum_{i=0}^p a_i \beta_i.$$

We estimate θ by $\hat{\theta} = \mathbf{a}^\top \hat{\boldsymbol{\beta}}$.

Example. Let $p = 2$ and let $\mathbf{a}^\top = (1, 1, 5)$. Then,

$$\theta = \mathbf{a}^\top \boldsymbol{\beta} = \beta_0 + \beta_1 + 5\beta_2.$$

Recall that $\hat{\boldsymbol{\beta}} \sim \text{MVN}(\boldsymbol{\beta}, \sigma^2(X^\top X)^{-1})$. Then,

$$\hat{\theta} = \mathbf{a}^\top \hat{\boldsymbol{\beta}} \sim \text{MVN}(\mathbf{a}^\top \boldsymbol{\beta}, \sigma^2 \mathbf{a}^\top (X^\top X)^{-1} \mathbf{a}).$$

We have,

$$\frac{\hat{\theta} - \theta}{\sigma \sqrt{\mathbf{a}^\top (X^\top X)^{-1} \mathbf{a}}} \sim N(0, 1), \quad \text{and} \quad \frac{\hat{\theta} - \theta}{\hat{\sigma} \sqrt{\mathbf{a}^\top (X^\top X)^{-1} \mathbf{a}}} \sim t_{n-p-1}.$$

Prediction Interval

Task: Predict y for given values of predictor variables $x_1, \dots, x_p$.

Let $\mathbf{a}_p = (1, x_1, \dots, x_p)^\top$. Since $y_i = \beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip} + \epsilon_i$, we have

$$y_p = \mathbf{a}_p^\top \boldsymbol{\beta} + \epsilon_p.$$

We plug in the LSE $\hat{\boldsymbol{\beta}}$ to estimate y_p :

$$\hat{y}_p = \mathbf{a}_p^\top \hat{\boldsymbol{\beta}}.$$

Note that the prediction error is $y_p - \hat{y}_p$. Now,

$$\begin{aligned}
 \text{Var}(y_p - \hat{y}_p) &= \text{Var}(\mathbf{a}_p^\top \boldsymbol{\beta} + \epsilon_p - \mathbf{a}_p^\top \hat{\boldsymbol{\beta}}) \\
 &= \text{Var}(\epsilon_p - \mathbf{a}_p^\top \hat{\boldsymbol{\beta}}) && \text{since } \mathbf{a}_p^\top \boldsymbol{\beta} \text{ is a constant} \\
 &= \text{Var}(\epsilon_p) + \text{Var}(\mathbf{a}_p^\top \hat{\boldsymbol{\beta}}) && \text{since } \epsilon_p \perp \hat{\boldsymbol{\beta}} \\
 &= \sigma^2 + \sigma^2 \mathbf{a}_p^\top (X^\top X)^{-1} \mathbf{a}_p \\
 &= \sigma^2 [1 + \mathbf{a}_p^\top (X^\top X)^{-1} \mathbf{a}_p].
 \end{aligned}$$


As usual, we plug in $\hat{\sigma}^2 = \frac{\text{SSE}}{n-p-1}$ in place of σ^2 . Thus,

$$\frac{y_p - \hat{y}_p}{\hat{\sigma} \sqrt{1 + \mathbf{a}_p^\top (X^\top X)^{-1} \mathbf{a}_p}} \sim t_{n-p-1}.$$

Proposition 2.9 (Prediction Interval for y_p).

A $100(1 - \alpha)\%$ prediction interval for y_p is

$$\hat{y}_p \pm t_{n-p-1, \alpha/2} \text{SE}(\hat{y}_p - y_p) \quad \text{where } \text{SE}(\hat{y}_p - y_p) = \hat{\sigma} \sqrt{1 + \mathbf{a}_p^\top (X^\top X)^{-1} \mathbf{a}_p}.$$



MLR in R

```
dat <- read.csv("lowbwt2.csv", header=T)
dim(dat)
head(dat)

pairs(dat)

library(car)
car::scatterplotMatrix(dat, col="black")

## Matrix form
X <- dat[,2:4] #  $Y = \beta X + \epsilon$ 
dim(X)

X <- as.matrix(cbind(1, X)) # appends a column of 1's in the front
dim(X)
head(X)

n <- nrow(X)
p <- ncol(X) - 1
n
p

Y <- dat$headcirc # equivalently, dat[,1]
Ybar <- mean(Y)

## LSE
beta <- solve(t(X) %*% X) %*% t(X) %*% Y #  $\beta = (X^t X)^{-1} X^t Y$ 
beta

## Hat matrix
```

```

H <- X %*% solve(t(X) %*% X) %*% t(X)    #  $H = X(X^t X)^{-1} X^t$ 
dim(H)
# Idempotent
max(abs(H %*% H - H))    #  $HH - H \approx 0$ , very close to 0
# symmetric
max(abs(t(H) - H))    #  $H^t - H \approx 0$ 
Yhat <- H %*% Y    #  $\hat{Y} = HY$ 

## Estimate  $\sigma^2$ 
resid <- Y-Yhat
SSE <- sum(resid^2)
sigma2 <- SSE/(n-p-1)
sigma <- sqrt(sigma2)
c(sigma, sigma2)

## Estimate variance-covariance matrix of  $\hat{\beta}$ 
XtX <- solve(t(X) %*% X)
XtX
V <- sigma2 * XtX
V

## Confidence interval for  $\beta_1$ 
beta[2] + c(-1, 1)*qt(0.975, n-p-1)* sqrt(V[2,2])

## Prediction interval
# for a child with gestage=30, birthwt=1200, toxemia=0
a <- c(1,30,1200,0)
a <- as.matrix(a, nrow=4, ncol=1)
# estimate
t(a) %*% beta

# variance
t(a) %*% V %*% a

```



```

# final prediction interval
c(t(a) %**% beta) + c(-1, 1)*qt(0.975, n-p-1)*
c(sqrt(1+t(a) %**% XtX %**% a))*sigma

### MLR using lm()
fm <- lm(headcirc~gestage+birthwt+toxemia, data=dat)
smy <- summary(fm)
names(smy)
coef(smy)

## Confidence interval for  $\beta_1$ 
smy$coefficients
smy$coefficients[2,2]
coef(fm)[2] + c(-1, 1)*qt(0.975, n-p-1)*smy$coefficients[2,2]
# Alternatively: direct method
confint(fm, 2)

## Estimate variance-covariance matrix of  $\hat{\beta}$ 
smy$cov.unscaled
smy$^sigma2 * smy$cov.unscaled

## Prediction interval
# for a child with gestage=30, birthwt=1200, toxemia=0
newdata <- data.frame(gestage=30, birthwt=1200, toxemia=0)
predict(fm, newdata)
predict(fm, newdata, interval="prediction", level=0.95)

```

2.5 MLR: ANOVA

MLR model:

$$y_i = \beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip} + \epsilon_i, \quad i = 1, \dots, n.$$

The sum of squares of residuals (error) is

$$\text{SSE}(\hat{\beta}) = \sum_{i=1}^n r_i^2 = \mathbf{r}^\top \mathbf{r} = (\mathbf{Y} - X\hat{\beta})^\top (\mathbf{Y} - X\hat{\beta}).$$

Test $H_0 : \beta_1 = \beta_2 = \cdots = \beta_p = 0$, recall that

$$\frac{\hat{\beta}_i - \beta_i}{\text{SE}(\hat{\beta}_i)} \sim t_{n-p-1}.$$

What if we reject H_0 if we reject any $H_{0i} : \beta_i = 0$ for $i = 1, \dots, p$ at level $\alpha = 0.05$?

Assume H_0 is true and $\hat{\beta}_i$'s are independent. Then,

$$\mathbb{P}(\text{reject } H_0 \mid H_0 \text{ is true}) = \mathbb{P}(\text{reject any } H_{0i} \mid H_0) = 1 - \mathbb{P}(\text{accept all } H_{0i} \mid H_0) = 1 - (1 - 0.05)^p.$$

Under H_0 , the full model is reduced to:

$$y_i = \beta_0 + \epsilon_i, \quad i = 1, \dots, n.$$

This is called the **reduced model** or **restricted model**. The LSE under the reduced model is

$$\hat{\beta}_0 = \bar{y},$$

and

$$\text{SSE}(\hat{\beta}_0) = \sum_{i=1}^n (y_i - \bar{y})^2 = \text{SST}.$$

Consider the following difference:

$$\text{SSE}(\hat{\beta}_0) - \text{SSE}(\hat{\beta}) = \text{SST} - \text{SSE}(\hat{\beta}) = \text{SSR}$$

where

$$\begin{aligned}
SSR &= \sum_{i=1}^n (y_i - \bar{y})^2 - (\mathbf{Y} - X\hat{\beta})^\top (\mathbf{Y} - X\hat{\beta}) \\
&= \sum_{i=1}^n y_i^2 - n\bar{y}^2 - [(I - H)\mathbf{Y}]^\top [(I - H)\mathbf{Y}] \\
&= \mathbf{Y}^\top \mathbf{Y} - n\bar{y}^2 - \mathbf{Y}^\top (I - H)^\top (I - H) \mathbf{Y} \\
&= \mathbf{Y}^\top \mathbf{Y} - n\bar{y}^2 - \mathbf{Y}^\top (I - H) \mathbf{Y} && \text{since } I - H \text{ is symmetric and idempotent} \\
&= \mathbf{Y}^\top \mathbf{Y} - n\bar{y}^2 - \mathbf{Y}^\top \mathbf{Y} + \mathbf{Y}^\top H \mathbf{Y} \\
&= \mathbf{Y}^\top H \mathbf{Y} - n\bar{y}^2 \\
&= \mathbf{Y}^\top H^\top H \mathbf{Y} - n\bar{y}^2 && \text{since } H \text{ is symmetric and idempotent} \\
&= \hat{\mathbf{Y}}^\top \hat{\mathbf{Y}} - n\bar{y}^2 \\
&= \hat{\beta}^\top X^\top X \hat{\beta} - n\bar{y}^2.
\end{aligned}$$

Proposition 2.10. Under $H_0 : \beta_1 = \dots = \beta_p = 0$, we have

$$\frac{SSR}{\sigma^2} \sim \chi_p^2.$$

Using this result, we can construct the F -statistic:

$$F = \frac{SSR/p}{SSE/(n-p-1)} = \frac{MSR}{MSE} \sim F_{p, n-p-1}.$$

We reject H_0 at level α if

$$F > F_\alpha(p, n-p-1), \quad \text{where } \alpha \text{ is the } 100(1-\alpha)\% \text{-tile.}$$

We have the following ANOVA Table:

Source of Variation	Sum of Squares	Degrees of Freedom	Mean Squares	F
Regression	$SSR = \hat{\beta}^T X^T X \hat{\beta} - n\bar{y}^2$	p	$MSR = \frac{SSR}{p}$	$\frac{MSR}{MSE}$
Residual	$SSE = (\mathbf{Y} - X\hat{\beta})^T (\mathbf{Y} - X\hat{\beta})$	$n - p - 1$	$MSE = \frac{SSE}{n - p - 1}$	
Total	$SST = \sum_{i=1}^n (y_i - \bar{y})^2$	$n - 1$		

Definition 2.12 (Total Coefficient of Determination).

The **total coefficient of determination** is

$$R^2 = \frac{SSR}{SST}$$

where $0 \leq R^2 \leq 1$.

Recall the MLR model is

$$y_i = \beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip} + \epsilon_i$$

and $y_i \sim N(\beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip}, \sigma^2)$ independently. The PDF of $y \sim N(\mu, \sigma^2)$ is

$$f(y) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y - \mu)^2}{2\sigma^2}\right).$$

The likelihood function is

$$L(\beta, \sigma^2) = \prod_{i=1}^n f(y_i) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{[y_i - (\beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip})]^2}{2\sigma^2}\right).$$

The log-likelihood function is

$$\begin{aligned}
 \ell(\boldsymbol{\beta}, \sigma^2) &= \log L(\boldsymbol{\beta}, \sigma^2) \\
 &= \sum_{i=1}^n \left(\log(2\pi\sigma^2)^{-1/2} - \frac{[y_i - (\beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip})]^2}{2\sigma^2} \right) \\
 &= -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \underbrace{\sum_{i=1}^n [y_i - (\beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip})]^2}_{S(\boldsymbol{\beta})}.
 \end{aligned}$$

For $\boldsymbol{\beta}$, maximizing $\ell(\boldsymbol{\beta}, \sigma^2) \iff$ minimizing $S(\boldsymbol{\beta})$. That is,

$$\hat{\boldsymbol{\beta}}_{\text{MLE}} = \hat{\boldsymbol{\beta}}_{\text{LSE}}, \quad \text{and} \quad S(\hat{\boldsymbol{\beta}}) = \text{SSE}.$$

For σ^2 ,

$$\frac{\partial \ell(\boldsymbol{\beta}, \sigma^2)}{\partial \sigma^2} = -\frac{n}{2} \cdot \frac{1}{2\pi\sigma^2} 2\pi + \frac{\text{SSE}}{2} \cdot \frac{1}{(\sigma^2)^2} = 0 \implies -n + \frac{\text{SSE}}{\sigma^2} = 0 \implies \hat{\sigma}_{\text{MLE}}^2 = \frac{\text{SSE}}{n}.$$

Lecture 12, 2025/10/21

2.6 Geometric Interpretation of LSE

Recall that $\mathbf{Y} = X\boldsymbol{\beta} + \boldsymbol{\epsilon}$ represents n simultaneous equations

$$y_i = \beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip} + \epsilon_i, \quad i = 1, \dots, n.$$

Note that $X = \begin{bmatrix} \vec{1} & \mathbf{x}_1 & \mathbf{x}_2 & \cdots & \mathbf{x}_p \end{bmatrix}$ where

$$\vec{1} = \begin{bmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix}, \quad \mathbf{x}_i = \begin{bmatrix} x_{i1} \\ x_{i2} \\ \vdots \\ x_{in} \end{bmatrix} \quad \text{for } i = 1, \dots, p.$$

The systematic part

$$\begin{aligned}
 X\boldsymbol{\beta} &= \begin{bmatrix} \vec{1} & \mathbf{x}_1 & \mathbf{x}_2 & \cdots & \mathbf{x}_p \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_p \end{bmatrix} \\
 &= \beta_0 \vec{1} + \beta_1 \mathbf{x}_1 + \beta_2 \mathbf{x}_2 + \cdots + \beta_p \mathbf{x}_p \\
 &= \sum_{i=0}^p \beta_i \mathbf{x}_i \quad \text{where } \mathbf{x}_0 = \vec{1}
 \end{aligned}$$

is a linear combination of the vectors $\mathbf{x}_i$.

Definition 2.13 (Column Space).

The **column space** of X is defined as

$$\begin{aligned}
 C(X) &= \{\text{all vectors that can be constructed as a linear combination of the columns of } X\} \\
 &= \{X\boldsymbol{\beta} : \boldsymbol{\beta} \in \mathbb{R}^{p+1}\} \\
 &= \{\beta_0 \vec{1} + \beta_1 \mathbf{x}_1 + \cdots + \beta_p \mathbf{x}_p : \beta_i \in \mathbb{R}, i = 0, 1, \dots, p\}.
 \end{aligned}$$

Note. $\hat{\mathbf{Y}} = X\hat{\boldsymbol{\beta}} \in C(X)$.


$C(X)$ spans a $(p + 1)$ -dimensional subspace inside the n -dimensional space (e.g. $\mathbb{R}^n$).

Note that

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

is a point in a n -dimensional space and represents a vector. The same is true for $\vec{1}, \mathbf{x}_1, \dots, \mathbf{x}_p$.

Example. $\beta_0 \mathbf{x}_0 + \beta_1 \mathbf{x}_1$ is a linear combination of two vectors $\mathbf{x}_0$ and $\mathbf{x}_1$, spanning a plane.



The LSE $\hat{\beta}$ minimizes


$$\text{SSE} = \sum_{i=1}^n r_i^2,$$

the length $\sqrt{r_1^2 + r_2^2 + \dots + r_n^2}$ of the residual vector


$$\mathbf{r} = \begin{bmatrix} r_1 \\ r_2 \\ \vdots \\ r_n \end{bmatrix}$$

subject to the constraint $\hat{\mathbf{Y}} \in C(X)$. Recall that $\mathbf{Y} = \hat{\mathbf{Y}} + \mathbf{r}$. The LSE tries to find $\hat{\mathbf{Y}} \in C(X)$ that is “closest” to $\mathbf{Y}$.

$P = 0$



$P = 1$



Note that $\hat{\mathbf{Y}} = H\mathbf{Y}$ is the **perpendicular projection** of $\mathbf{Y}$ onto $C(X)$.

Since $\mathbf{r} \perp C(X)$, or equivalently,

$$\begin{aligned}\vec{1}^\top \mathbf{r} &= 0, \\ \mathbf{x}_1^\top \mathbf{r} &= 0, \\ &\vdots \\ \mathbf{x}_p^\top \mathbf{r} &= 0.\end{aligned}$$

(inner product = 0). This implies that

$$\begin{aligned}X^\top \mathbf{r} &= \vec{0} \\ X^\top (\mathbf{Y} - X\hat{\beta}) &= \vec{0} \\ X^\top X\hat{\beta} &= X^\top \mathbf{Y} \\ \hat{\beta} &= (X^\top X)^{-1} X^\top \mathbf{Y}.\end{aligned}$$

Proposition 2.11.

(1) $\mathbf{Y}$ is projected onto two orthogonal subspaces $C(X)$ and $C(X)^\perp$:

$$\mathbf{Y} = \hat{\mathbf{Y}} + \mathbf{r}, \quad \hat{\mathbf{Y}} \in C(X), \quad \mathbf{r} \in C(X)^\perp$$

where $C(X)$ has dimension $p + 1$ and $C(X)^\perp$ has dimension $n - p - 1$.

(2) Under the normality assumption, $\hat{\mathbf{Y}}$ and $\mathbf{r}$ are independent, i.e. $\hat{\mathbf{Y}} \perp\!\!\!\perp \mathbf{r}$.

Test Linear Constraints

Example. Test $H_0 : \beta_q = \beta_{q+1} = \dots = \beta_p = 0$ where $1 \leq q \leq p$. We can write H_0 in the form of

$$A\beta = \vec{0}$$

where

$$A = \begin{bmatrix} 0 & 0 & \cdots & 0 & 1 & 0 & \cdots & 0 \\ 0 & 0 & \cdots & 0 & 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 0 & 0 & 0 & \cdots & 1 \end{bmatrix}_{(p-q+1) \times (p+1)}.$$

The first “1” entry in the first row is in the $(q + 1)$ th column.

Example. Consider

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \epsilon$$

and the literature suggests

$$\beta_1 = 2\beta_2, \quad \text{and} \quad \beta_3 = 0.$$

We can write this in matrix form $A\boldsymbol{\beta} = \vec{0}$ where

$$A = \begin{bmatrix} 0 & 1 & -2 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}.$$

Suppose A is a $\ell \times (p + 1)$ matrix, i.e. we have ℓ linear constraints. Our test procedure relies on the **additional sum of squares principle**. Recall that

$$C(X) = \{\beta_0 \vec{1} + \beta_1 \mathbf{x}_1 + \cdots + \beta_p \mathbf{x}_p\}$$

We define


$$C_A(X) = \{\beta_0 \vec{1} + \beta_1 \mathbf{x}_1 + \cdots + \beta_p \mathbf{x}_p : A\boldsymbol{\beta} = \vec{0}\} \subseteq C(X)$$

which is a subspace of $C(X)$.

Example. For

$$\beta_1 \begin{bmatrix} 1 \\ 0 \end{bmatrix} + \beta_2 \begin{bmatrix} 0 \\ 1 \end{bmatrix} = \begin{bmatrix} \beta_1 \\ \beta_2 \end{bmatrix}.$$

The linear constraint is $\beta_1 = \beta_2$.




Let $\hat{\mathbf{Y}}$ be the orthogonal projection of $\mathbf{Y}$ onto $C(X)$, and let $\hat{\mathbf{Y}}_A$ be the orthogonal projection of $\mathbf{Y}$ onto $C_A(X)$. If $H_0 : A\boldsymbol{\beta} = \vec{0}$ is true, we expect $\hat{\mathbf{Y}}$ and $\hat{\mathbf{Y}}_A$ to be similar/close. That is, the model makes similar predictions whether we set $A\boldsymbol{\beta} = \vec{0}$ or not when fitting the model.

The squared distance between $\hat{\mathbf{Y}}$ and $\hat{\mathbf{Y}}_A$

$$\|\hat{\mathbf{Y}} - \hat{\mathbf{Y}}_A\|^2 = (\hat{\mathbf{Y}} - \hat{\mathbf{Y}}_A)^\top (\hat{\mathbf{Y}} - \hat{\mathbf{Y}}_A)$$

is the **additional sum of squares**.

Geometric representation of $C(X)$ and $C_A(X)$:



By Pythagoras' theorem,

$$(AC)^2 = \text{SSE}_A, \quad (BC)^2 = \text{SSE}$$

The squared distance is

$$\|\hat{\mathbf{Y}} - \hat{\mathbf{Y}}_A\|^2 = (AB)^2 = (AC)^2 - (BC)^2 = \text{SSE}_A - \text{SSE}$$

where $\text{SSE}_A = \|\mathbf{Y} - \hat{\mathbf{Y}}_A\|^2$ and $\text{SSE} = \|\mathbf{Y} - \hat{\mathbf{Y}}\|^2$.

Theorem 2.12. Under $H_0 : A\boldsymbol{\beta} = \vec{0}$ where A is a $\ell \times (p+1)$ matrix, we have

(1) $\frac{\|\hat{\mathbf{Y}} - \hat{\mathbf{Y}}_A\|^2}{\sigma^2} \sim \chi_\ell^2$.

(2) $\|\mathbf{Y} - \hat{\mathbf{Y}}\|^2$ is independent of

$$\hat{\sigma}^2 = \frac{(\mathbf{Y} - \hat{\mathbf{Y}})^\top (\mathbf{Y} - \hat{\mathbf{Y}})}{n - p - 1}.$$

That is, $\|\hat{\mathbf{Y}} - \hat{\mathbf{Y}}_A\|^2 \perp \hat{\sigma}^2$.

Lecture 13, 2025/10/23

The theorem above implies the following F -statistic (under $H_0 : A\boldsymbol{\beta} = \vec{0}$):

$$F = \frac{\|\hat{\mathbf{Y}} - \hat{\mathbf{Y}}_A\|^2 / \ell}{\hat{\sigma}^2} \sim F(\ell, n - p - 1).$$

Note. We saw earlier that the ratio of two independent chi-square random variables leads to an F -distribution. Large values of F provide evidence against H_0 .

We reject H_0 at level α if

$$F > F_\alpha(\ell, n - p - 1),$$

or we compute the p -value as

$$p\text{-value} = \mathbb{P}(F(\ell, n - p - 1) > F).$$

In summary, to test $H_0 : A\boldsymbol{\beta} = \vec{0}$, we follow these steps:

- (1) Fit the full model without constraints.
- (2) Compute SSE.
- (3) Fit the restricted model with constraints $A\beta = \vec{0}$.
- (4) Compute SSE_A , the sum of squares residuals under the restricted model.
- (5) Compute the F -statistic:

$$F = \frac{(SSE_A - SSE)/\ell}{SSE/(n - p - 1)}.$$

Question: How to fit the restricted model with constraints $A\beta = \vec{0}$?

Example (Continued).

Back to Example 2.6, we have

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \epsilon$$

with $\beta_1 = 2\beta_2$ and $\beta_3 = 0$. Thus,

$$y = \beta_0 + 2\beta_2 x_1 + \beta_2 x_2 + \epsilon = \beta_0 + \beta_2(2x_1 + x_2) + \epsilon.$$

Case Study: UFFI Example

Lecture 14, 2025/10/28

UFFI example continued.

Model 2: $\mu = \beta_0 + \beta_1 x_1 + \beta_2 x_2$.

...

Model 3: $\mu = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3$.

UFFI	Air	$\mathbf{x}$	$\mu = \mathbf{x}^T \boldsymbol{\beta}$
Present	$x + 1$	$\begin{pmatrix} 1 & 1 & x + 1 & x + 1 \end{pmatrix}$	$\beta_0 + \beta_1 + (\beta_2 + \beta_3)(x + 1)$
Present	x	$\begin{pmatrix} 1 & 1 & x & x \end{pmatrix}$	$\beta_0 + \beta_1 + (\beta_2 + \beta_3)x$
Absent	$x + 1$	$\begin{pmatrix} 1 & 0 & x + 1 & 0 \end{pmatrix}$	$\beta_0 + \beta_2(x + 1)$
Absent	x	$\begin{pmatrix} 1 & 0 & x & 0 \end{pmatrix}$	$\beta_0 + \beta_2 x$

Explaining the coefficients:

- Row 2 – Row 4 = $\beta_1 + \beta_3$ = expected increase of CH_2O level when UFFI is absent $\rightarrow$ present at air level x .
- Row 3 – Row 4 = β_2 = expected increase of CH_2O level when air level increases by 1 unit with UFFI absent.
- Row 1 – Row 2 = $\beta_2 + \beta_3$ = expected increase of CH_2O level when air level increases by 1 unit with UFFI present.

...

Interpretation: The effect of one predictor on the response depends on the levels of other predictors.

...

3 Model Specification

In MLR, we have

$$\mathbf{Y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon} \implies \mathbb{E}[\mathbf{Y}] = \mathbf{X}\boldsymbol{\beta}.$$

Suppose we have (x_i, y_i) for $i = 1, 2, \dots, n$.

plot...

We can extend to a quadratic model (also linear):

$$\mathbb{E}[y] = \beta_0 + \beta_1 x + \beta_2 x^2.$$


In matrix form:

$$\mathbb{E}[\mathbf{Y}] = \mathbf{X}\boldsymbol{\beta}.$$

Equivalently,

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \vdots & \vdots & \vdots \\ 1 & x_n & x_n^2 \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \end{bmatrix} + \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_n \end{bmatrix}.$$

Suppose we have a piecewise constant model:



Question: Can we make this a linear model?

Answer: Yes. Let's define an indicator variable:

$$x_1 = \mathbb{1}_{x \geq a} = I(x \geq a) = \begin{cases} 1 & x \geq a, \\ 0 & x < a. \end{cases}$$

Then, $y = \beta_0 + \beta_1 x_1$. So we have

$$X = \begin{bmatrix} 1 & I(x_1 \geq a) \\ 1 & I(x_2 \geq a) \\ \vdots & \vdots \\ 1 & I(x_n \geq a) \end{bmatrix} \quad \text{and} \quad \boldsymbol{\beta} = \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}.$$

To approximate a more complicated function, we can use many piecewise constant segments.

Now, suppose we have a piecewise linear model:

plot...

Then, we have

$$\begin{cases} y = \beta_0 + \beta_1 x & x < a, \\ y = \beta_2 + \beta_3 x & x \geq a. \end{cases}$$

We can rewrite this as


$$y = \beta_0 I(x < a) + \beta_1 x I(x < a) + \beta_2 I(x \geq a) + \beta_3 x I(x \geq a).$$

Thus, we have

$$X = \begin{bmatrix} I(x_1 < a) & x_1 I(x_1 < a) & I(x_1 \geq a) & x_1 I(x_1 \geq a) \\ I(x_2 < a) & x_2 I(x_2 < a) & I(x_2 \geq a) & x_2 I(x_2 \geq a) \\ \vdots & \vdots & \vdots & \vdots \\ I(x_n < a) & x_n I(x_n < a) & I(x_n \geq a) & x_n I(x_n \geq a) \end{bmatrix} \quad \text{and} \quad \boldsymbol{\beta} = \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \beta_3 \end{bmatrix}.$$

Then, suppose we have a piecewise linear model with continuity:

plot...



The model is:

$$\begin{cases} \mathbb{E}[y] = \beta_0 + \beta_1 x & x < a, \\ \mathbb{E}[y] = \beta_2 + \beta_3 x & x \geq a \end{cases}$$

with the constraint:

$$\beta_0 + \beta_1 a = \beta_2 + \beta_3 a \iff \beta_2 = \beta_0 + (\beta_1 - \beta_3)a.$$

Then, the model becomes:

$$\begin{cases} \mathbb{E}[y] = \beta_0 + \beta_1 x & x < a, \\ \mathbb{E}[y] = \underbrace{\beta_0 + (\beta_1 - \beta_3)a}_{\beta_2} + \beta_3 x \\ = \beta_0 + \beta_1 a + \beta_3(x - a) & x \geq a. \end{cases}$$

We can use the indicator function to further simplify:

$$\begin{aligned} \mathbb{E}[y] &= \beta_0 I(x < a) + \beta_1 x I(x < a) + (\beta_0 + \beta_1 a) I(x \geq a) + \beta_3(x - a) I(x \geq a) \\ &= \beta_0 + \underbrace{\beta_1 x - \beta_1 x I(x \geq a)}_{\beta_1 x I(x < a)} + \beta_1 a I(x \geq a) + \beta_3(x - a) I(x \geq a) \\ &= \beta_0 + \beta_1 x + \underbrace{(\beta_3 - \beta_1)(x - a)}_{\beta_4} I(x \geq a). \end{aligned}$$

Question: What if we have several break points? (In the plot above, only have one at $x = a$.)

Question: If we know a , we can fit a linear model, but what if we don't know a ?

Answer: Try a that minimizes SSE.

One Sample Problem:

$$y_i = \beta_0 + \epsilon_i \quad \text{with} \quad \mathbb{E}[y_i] = \beta_0.$$

Then,

$$Y = X\beta \quad \text{with} \quad Y = \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix}, \quad X = \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix}, \quad \beta = \beta_0.$$

Two-Sample Problem: “Cell means” model:

$$y_{ij} = \mu_i + \epsilon_{ij} \quad i = 1, 2.$$

Suppose that group i has n_i observations. Then,

$$\mathbb{E}[Y] = X\beta \quad \text{with}$$

$$\mathbb{E}[Y] = \mathbb{E} \begin{bmatrix} y_{11} \\ \vdots \\ y_{1n_1} \\ y_{21} \\ \vdots \\ y_{2n_2} \end{bmatrix} = \begin{bmatrix} \mu_1 \\ \vdots \\ \mu_1 \\ \mu_2 \\ \vdots \\ \mu_2 \end{bmatrix} = \begin{bmatrix} 1 \\ \vdots \\ 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix} \mu_1 + \begin{bmatrix} 0 \\ \vdots \\ 0 \\ 1 \\ \vdots \\ 1 \end{bmatrix} \mu_2 = \underbrace{\begin{bmatrix} \mathbf{1}_{n_1} & \mathbf{0}_{n_1} \\ \mathbf{0}_{n_2} & \mathbf{1}_{n_2} \end{bmatrix}}_X \underbrace{\begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix}}_\beta.$$

Exercise: What is $X^\top X$?

Answer: we have

$$\begin{aligned}
 X^T X &= \begin{bmatrix} \mathbf{1}_{n_1}^T & \mathbf{0}_{n_2}^T \\ \mathbf{0}_{n_1}^T & \mathbf{1}_{n_2}^T \end{bmatrix}_{2 \times n} \begin{bmatrix} \mathbf{1}_{n_1} & \mathbf{0}_{n_1} \\ \mathbf{0}_{n_2} & \mathbf{1}_{n_2} \end{bmatrix}_{n \times 2} \\
 &= \begin{bmatrix} 1 & \dots & 1 & 0 & \dots & 0 \\ 0 & \dots & 0 & 1 & \dots & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ \vdots & \vdots \\ 1 & 0 \\ 0 & 1 \\ \vdots & \vdots \\ 0 & 1 \end{bmatrix} = \begin{bmatrix} n_1 & 0 \\ 0 & n_2 \end{bmatrix}.
 \end{aligned}$$

Exercise: What is $X^T Y$?

Answer: we have

$$X^T Y = X^T X \beta = \begin{bmatrix} n_1 & 0 \\ 0 & n_2 \end{bmatrix} \begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix} = \begin{bmatrix} n_1 \mu_1 \\ n_2 \mu_2 \end{bmatrix}.$$

We can also compute directly:

$$X^T Y = \begin{bmatrix} \mathbf{1}_{n_1}^T & \mathbf{0}_{n_2}^T \\ \mathbf{0}_{n_1}^T & \mathbf{1}_{n_2}^T \end{bmatrix} \begin{bmatrix} y_{11} \\ \vdots \\ y_{1n_1} \\ y_{21} \\ \vdots \\ y_{2n_2} \end{bmatrix} = \begin{bmatrix} \sum_{j=1}^{n_1} y_{1j} \\ \sum_{j=1}^{n_2} y_{2j} \end{bmatrix}.$$

Also, we have

$$\hat{\beta} = \begin{bmatrix} \hat{\mu}_1 \\ \hat{\mu}_2 \end{bmatrix} = (X^T X)^{-1} X^T Y = \begin{bmatrix} \frac{1}{n_1} & 0 \\ 0 & \frac{1}{n_2} \end{bmatrix} \begin{bmatrix} \sum_{j=1}^{n_1} y_{1j} \\ \sum_{j=1}^{n_2} y_{2j} \end{bmatrix} = \begin{bmatrix} \bar{y}_{1+} \\ \bar{y}_{2+} \end{bmatrix}$$

where $\bar{y}_{i+} = \hat{\mu}_i = \frac{\sum_{j=1}^{n_i} y_{ij}}{n_i}$.

Effects Model:

$$\mathbb{E}[y_i] = \beta_0 + \beta_1 x_i \quad \text{where} \quad x_i = I(\text{observation } i \text{ in group 2}).$$

Then,

$$\mathbb{E}[Y] = \mathbb{E} \begin{bmatrix} y_{11} \\ \vdots \\ y_{1n_1} \\ y_{21} \\ \vdots \\ y_{2n_2} \end{bmatrix} = \underbrace{\begin{bmatrix} \mathbf{1}_{n \times 1} & \begin{bmatrix} \mathbf{0}_{n_1 \times 1} \\ \mathbf{1}_{n_2 \times 1} \end{bmatrix} \end{bmatrix}}_X \underbrace{\begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}}_{\beta}.$$

Exercises: What are $X^\top X$, $X^\top Y$, and $\hat{\beta}$?

Answers: We have

$$X^\top X = \begin{bmatrix} 1 & \cdots & \cdots & \cdots & \cdots & 1 \\ 0 & \cdots & 0 & 1 & \cdots & 1 \end{bmatrix}_{2 \times n} \begin{bmatrix} 1 & 0 \\ \vdots & \vdots \\ 1 & 0 \\ 1 & 1 \\ \vdots & \vdots \\ 1 & 1 \end{bmatrix}_{n \times 2} = \begin{bmatrix} n & n_2 \\ n_2 & n_2 \end{bmatrix}$$

$$X^\top Y = \begin{bmatrix} 1 & \cdots & \cdots & \cdots & \cdots & 1 \\ 0 & \cdots & 0 & 1 & \cdots & 1 \end{bmatrix} \begin{bmatrix} y_{11} \\ \vdots \\ y_{1n_1} \\ y_{21} \\ \vdots \\ y_{2n_2} \end{bmatrix} = \begin{bmatrix} \sum_{i,j} y_{ij} \\ \sum_{j=1}^{n_2} y_{2j} \end{bmatrix}$$

Also,

$$\hat{\beta} = \begin{bmatrix} \hat{\beta}_0 \\ \hat{\beta}_1 \end{bmatrix} = (X^\top X)^{-1} X^\top Y = \begin{bmatrix} \bar{y}_{1+} \\ \bar{y}_{2+} - \bar{y}_{1+} \end{bmatrix}.$$

***k*-Sample Problem:**

Data:

Group	Observations
1	$y_{11}, y_{12}, \dots, y_{1n_1}$
2	$y_{21}, y_{22}, \dots, y_{2n_2}$
$\vdots$	$\vdots$
k	$y_{k1}, y_{k2}, \dots, y_{kn_k}$

The “cell means” model:

$$y_{ij} = \mu_i + \epsilon_{ij} \quad i = 1, 2, \dots, k.$$

Then.

$$\mathbb{E} \begin{bmatrix} y_{11} \\ \vdots \\ y_{1n_1} \\ y_{k1} \\ \vdots \\ y_{kn_k} \end{bmatrix} = \begin{bmatrix} \mathbf{1}_{n_1 \times 1} & \mathbf{0} & \cdots & \mathbf{0} \\ \mathbf{0} & \mathbf{1}_{n_2 \times 1} & \cdots & \mathbf{0} \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{0} & \mathbf{0} & \cdots & \mathbf{1}_{n_k \times 1} \end{bmatrix} \begin{bmatrix} \mu_1 \\ \mu_2 \\ \vdots \\ \mu_k \end{bmatrix}.$$

We have

$$X^\top X = \begin{bmatrix} n_1 & 0 & \cdots & 0 \\ 0 & n_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & n_k \end{bmatrix} = \text{diag}(n_1, n_2, \dots, n_k) \quad X^\top Y = \begin{bmatrix} \sum_{j=1}^{n_1} y_{1j} \\ \sum_{j=1}^{n_2} y_{2j} \\ \vdots \\ \sum_{j=1}^{n_k} y_{kj} \end{bmatrix}.$$

Also,

$$\hat{\beta} = (X^\top X)^{-1} X^\top Y = \begin{bmatrix} \bar{y}_{1+} \\ \bar{y}_{2+} \\ \vdots \\ \bar{y}_{k+} \end{bmatrix} \quad \text{where} \quad \bar{y}_{i+} = \frac{\sum_{j=1}^{n_i} y_{ij}}{n_i}.$$

And,

$$\hat{y} = X\hat{\beta} = \left[\begin{array}{c} \bar{y}_{1+} \\ \vdots \\ \bar{y}_{1+} \\ \vdots \\ \bar{y}_{k+} \\ \vdots \\ \bar{y}_{k+} \end{array} \right] \left\{ \begin{array}{l} n_1 \\ \\ n_k \end{array} \right.$$

Effects Model:

$$\mathbb{E}[y_i] = \beta_1 + \beta_2 x_{i2} + \cdots + \beta_k x_{ik} \quad \text{where} \quad x_{ij} = I(\text{observation } i \text{ in group } j)$$

and

$$\mathbb{E}[Y] = \underbrace{\begin{bmatrix} \mathbf{1}_{n \times 1} & \mathbf{x}_2 & \cdots & \mathbf{x}_k \end{bmatrix}}_X \underbrace{\begin{bmatrix} \beta_1 \\ \beta_2 \\ \vdots \\ \beta_k \end{bmatrix}}_{\beta} \quad \text{where} \quad \mathbf{x}_j = \begin{bmatrix} x_{1j} \\ x_{2j} \\ \vdots \\ x_{nj} \end{bmatrix}.$$

Also,

$$\hat{\beta} = (X^T X)^{-1} X^T Y = \begin{bmatrix} \bar{y}_{1+} \\ \bar{y}_{2+} - \bar{y}_{1+} \\ \vdots \\ \bar{y}_{k+} - \bar{y}_{1+} \end{bmatrix} \quad \text{where } \bar{y}_{1+} \text{ is the reference group.}$$

ANOVA

We have

$$\begin{aligned}
 SSR &= \hat{\beta}^\top X^\top X \hat{\beta} - n\bar{y}^2 \\
 &= \hat{\beta}^\top X^\top Y - n\bar{y}^2 \\
 &= \begin{bmatrix} \bar{y}_{1+} & \bar{y}_{2+} & \cdots & \bar{y}_{k+} \end{bmatrix} \begin{bmatrix} n_1 & \bar{y}_{1+} \\ \vdots & \vdots \\ n_k & \bar{y}_{k+} \end{bmatrix} - n\bar{y}^2 \quad \text{using cell means results} \\
 &= \sum_{i=1}^k n_i \bar{y}_{i+}^2 - n\bar{y}^2 \\
 &= \sum_{i=1}^k n_i (\bar{y}_{i+} - \bar{y})^2.
 \end{aligned}$$

Also,

$$SSE = (Y - \hat{Y})^\top (Y - \hat{Y}) = \sum_{i=1}^k \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{i+})^2.$$

ANOVA Table

Source	Sum of Squares	Degrees of Freedom	Mean Squares	F
Regression	$SSR = \sum_{i=1}^k n_i (\bar{y}_{i+} - \bar{y})^2$	$k - 1$	$MSR = \frac{SSR}{k - 1}$	$\frac{MSR}{MSE}$
Residual	$SSE = \sum_{i=1}^k \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{i+})^2$	$n - k$	$MSE = \frac{SSE}{n - k}$	
Total	$SST = \sum_{i=1}^n (y_i - \bar{y})^2$	$n - 1$		

4 Model Checking

Lecture 17, 2025/11/06

Residuals

- Residual: the difference between the observed y and the fitted y : $r_i = y_i - \hat{y}_i$.
- Random error: the difference between the observed y and the expected y : $\epsilon_i = y_i - \mathbb{E}[y_i]$.

Proposition 4.1. $\mathbf{r} = (I - H)\epsilon$.

Proof. We have

$$\begin{aligned}\mathbf{r} &= Y - \hat{Y} \\ &= Y - HY \\ &= (I - H)Y \\ &= (I - H)(X\beta + \epsilon).\end{aligned}$$

It is enough to show that $(I - H)X\beta = \vec{0}$:

$$X\beta - X(X^\top X)^{-1}X^\top X\beta = X\beta - X\beta = \vec{0}.$$

Note that H is idempotent, i.e. $H = HH$, then

$$h_{ii} = (H)_{ii} = (HH)_{ii} = \sum_{j=1}^n h_{ij}h_{ji}.$$

Also, H is symmetric, i.e. $h_{ij} = h_{ji}$, so

$$h_{ii} = \sum_{j=1}^n h_{ij}^2 = h_{ii}^2 + \sum_{j \neq i} h_{ij}^2.$$

Then,

$$\sum_{j \neq i} h_{ij}^2 = h_{ii} - h_{ii}^2 = h_{ii}(1 - h_{ii}) \geq 0.$$

We also know that $0 \leq h_{ii} \leq 1$ since if $h_{ii} > 1$, then $1 - h_{ii} < 0$ which is a contradiction. Thus,

$$0 \leq h_{ii}(1 - h_{ii}) \leq \frac{1}{4} \quad \text{by setting derivative to 0.}$$

That is, $\sum_{j \neq i} h_{ij}^2 \leq \frac{1}{4}$. The conclusion is that the off-diagonal elements cannot be large. Note that

$$\sum_{j=1}^n h_{ij} = \text{tr}(H) = \text{tr}(X(X^\top X)^{-1}X^\top) = \text{tr}((X^\top X)^{-1}X^\top X) = \text{tr}(I_{p+1}) = p + 1.$$

The average of h_{ii} 's is $\frac{p+1}{n}$. If $n \gg p$, all elements of H are expected to be small, thus, $\mathbf{r} = (I - H)\boldsymbol{\epsilon}$. $\square$

Exercise: What if $n = p + 1$? What would be H and $\mathbf{r}$?

Proposition 4.2. $\text{Cov}(X, \mathbf{r}) = \vec{0}$.

Proof. We have

$$\text{Cov}(X, \mathbf{r}) = \text{Cov}((I - H)Y, HY) = (I - H)\sigma^2 IH = \sigma^2(H - HH) = \vec{0}.$$

$\square$